

The Robobatter-9000 Strikes Out!

Using Machine Learning to Classify and Predict MLB Pitches

Introduction

The year is 2123. Self-learning robots are now common. One robot in question is officially named the "ROBOBATTER-9000", but goes by Bob after realizing that there is a critical lack of sports players who go by Bob. In this technologically near land, Bob wants nothing more than to take over the world... of baseball.

Because Bob is a specialized robot, he has mastered his swing to be mechanically and has advanced vision and telemetry can track almost any pitch. If he can connect with the ball, the ball will be crushed for a home run. However, a design flaw has taken away almost all of the pre-pitch time away to set up the mechanics of the swing and the telemetry. Because of that, Bob will have to decide pre-pitch what type of pitch will be thrown. If it is a fastball, he can easily track it and hit it out of the park. Otherwise, he will pass on the pitch if not a fastball, and will wait until another fastball is thrown to take another chance at swinging. This leaves Bob with 2 main problems: 1) Can he correctly identify and classify pitches based on who is throwing them? 2) Can he correctly predict what type of pitch is coming next?

Part 1: Pitch Classification

The Dataset

The main dataset used will be the pybaseball package, which scrapes data from Baseball Reference, MLB Statcast, and Fangraphs. The primary utilization will be the Statcast portion of the dataset, used for its detailed pitch breakdowns, such as pitch x, y, and z velocity and acceleration.

```
In [ ]: !pip install pybaseball
```

```
import pybaseball as pyb
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.model_selection import train_test_split
from sklearn import tree
from sklearn.ensemble import RandomForestClassifier
from sklearn import metrics
from sklearn.metrics import accuracy_score
from sklearn.metrics import classification_report

import warnings
warnings.filterwarnings('ignore')
```

Pybaseball has a statcast function, which will gather every single pitch thrown in the 2022 MLB season, which is over 750,000 pitches.

```
In [ ]: pitches = pyb.statcast(start_dt='2022-03-22', end_dt='2022-10-05')
print("Dataset has {} pitches".format(len(pitches)))
print(pitches.shape)
```

This is a large query, it may take a moment to complete
100%|██████████| 198/198 [01:59<00:00, 1.66it/s]
Dataset has 750075 pitches
(750075, 92)

First we will begin to clean up the data. The first part of cleaning up the data will be identifying and getting rid of null items in the dataset.

```
In [ ]: pitches.isnull().sum()
```

```
Out[ ]: pitch_type      19901
game_date              0
release_speed         19666
release_pos_x         19649
release_pos_z         19649
...
if_fielding_alignment  30174
of_fielding_alignment  30174
spin_axis             22751
delta_home_win_exp     0
delta_run_exp         41582
Length: 92, dtype: int64
```

```
In [ ]: pitches.dropna(axis="columns", how="all", inplace=True)
pitches.shape
```

```
Out[ ]: (750075, 84)
```

The next step in cleaning up the dataset is identifying what types of pitches we want to include. Because Statcast identifies pitches like pitchouts (when a pitcher throws the ball to the catcher so that the catcher can throw out a baserunner that is stealing a base) or intentional balls (when a pitcher intentionally throws a ball to eventually walk the batter), we need to filter out to only include actual pitches thrown to better categorize the dataset.

```
In [ ]: pitches.pitch_type.value_counts()
```

```
Out[ ]: FF      242421
        SL      132022
        SI      113493
        CH      82708
        CU      57038
        FC      52043
        ST      19694
        KC      15493
        FS      10986
        SV       2319
        FA       1279
        EP        494
        CS        108
        PO         40
        KN         19
        SC         17
        Name: pitch_type, dtype: int64
```

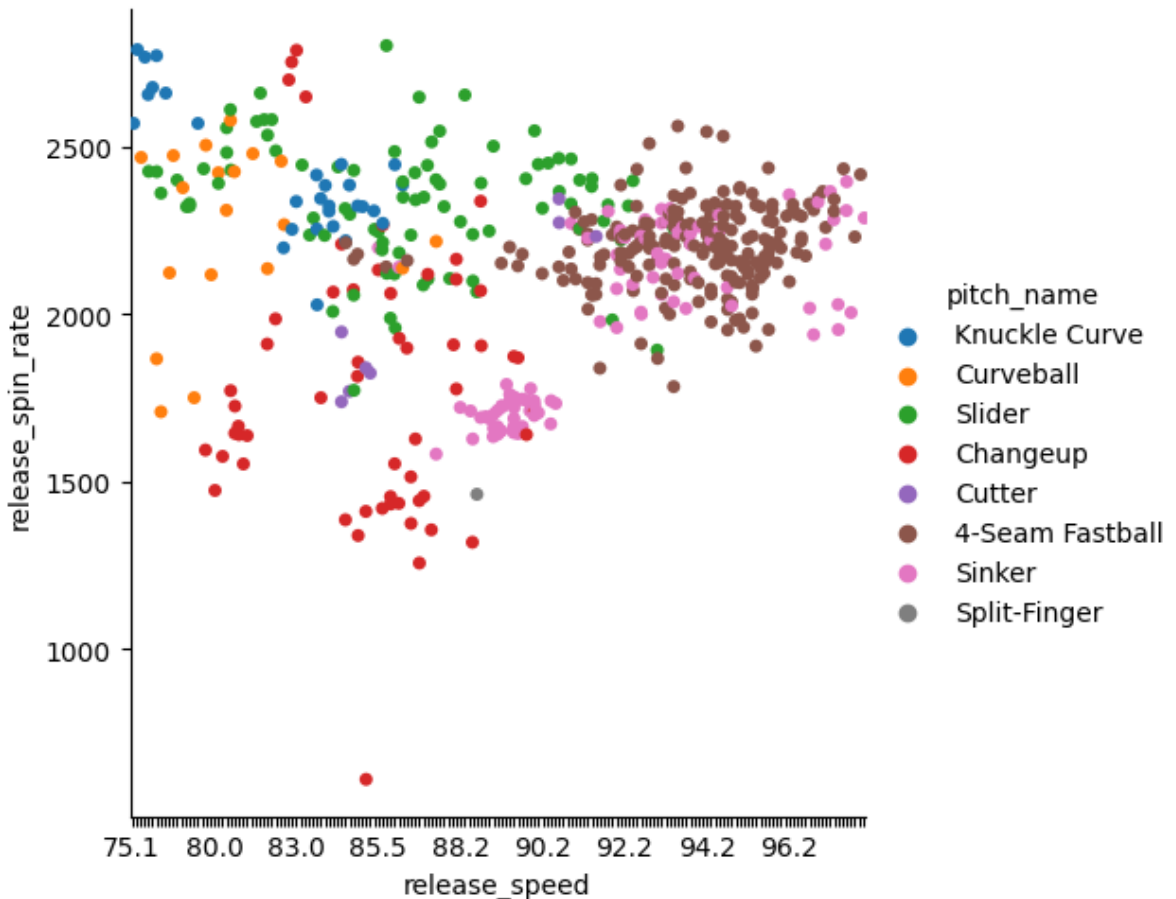
```
In [ ]: pitches=pitches[pitches.pitch_type.isin(['CH', 'CU', 'EP', 'FC', 'FF', 'FO', 'FS'],
        pitches.pitch_type.value_counts())
```

```
Out[ ]: FF      242421
        SL      132022
        SI      113493
        CH      82708
        CU      57038
        FC      52043
        KC      15493
        FS      10986
        EP        494
        KN         19
        SC         17
        Name: pitch_type, dtype: int64
```

To get a sense of what types of pitches are being thrown, we will take the first 500 pitches of the dataset and categorize them based on release speed and release spin rate. When graphed, we can get a sense of "clusters" of types of pitches. The fastballs shown (4-seam, sinker) are grouped very close together. Pretty much everything else can be classified as an "offspeed" pitch, and are each individually clustered on its own, not clustered with other types of offspeed pitches.

```
In [ ]: plot = sns.catplot(x='release_speed', y = 'release_spin_rate', hue = 'pitch_name',

for ind, label in enumerate(plot.ax.get_xticklabels()):
    if ind % 20 == 0:
        label.set_visible(True)
    else:
        label.set_visible(False)
```



To help with the model later, we will create some extra variables, which will if there are runners on first, second, or third, and the total number of baserunners. For the pitch classification step next, the pitch will be either labeled as a fastball, breaking ball, or change. Breaking down just the fastballs, the dataset is roughly 60% fastballs and 40% offspeed (which includes both the breaking balls and changes).

```
In [ ]: pitches['runner_on_first'] = np.where(~pitches.on_1b.isnull(), 1, 0)
pitches['runner_on_second'] = np.where(~pitches.on_2b.isnull(), 1, 0)
pitches['runner_on_third'] = np.where(~pitches.on_3b.isnull(), 1, 0)
pitches['total_baserunners'] = pitches['runner_on_first'] + pitches['runner_on_sec
pitches['is_fastball'] = np.where(pitches.pitch_type.isin(['FF', 'SI', 'FT', 'FA', 'FS'
pitches['is_breaking'] = np.where(pitches.pitch_type.isin(['CU', 'KC', 'KN', 'SC',
pitches['is_change'] = np.where(pitches.pitch_type.isin(['CH', 'EP']), 1, 0)

pitches.is_fastball.value_counts(normalize=True)
```

```
Out[ ]: 1    0.592787
0    0.407213
Name: is_fastball, dtype: float64
```

Since we want this classification model to run on several different datasets, we can create a function that cleans, trains, and tests the model in one action. We will first only use a subset of the dataset, utilizing the pitch name, release spin and spin rate, as well as the x, y, and z pitch velocity and acceleration. We will also create "filters" for the fastball, breaking ball, and change pitches to see if the model can classify the pitches in grouped and non-grouped scenarios. We will use a decision tree as the primary machine learning model, and then compare the training and test score accuracies as well as the overall classification report of the decision tree.

```
In [ ]: def decision_tree(data, fastball, breaking, change):
    data = data.loc[:, ['pitch_name',
                        'release_speed',
                        'release_spin_rate',
                        'vx0', 'vy0', 'vz0',
                        'ax', 'ay', 'az']]

    data = data.dropna()

    def fastball_pitch_filter(x):
        if x=='2-Seam Fastball' or x=='4-Seam Fastball' or x=='Sinker' or x=='Split':
            return 'Fastball'
        return x

    def breaking_pitch_filter(x):
        if x=='Curveball' or x=='Slider' or x=='Screwball' or x=='Knuckle Curve' or x=='Changeup':
            return 'Breaking'
        return x

    def change_pitch_filter(x):
        if x=='Changeup' or x=='Eephus':
            return 'Change'
        return x

    if fastball == True:
        data['pitch_name'] = data['pitch_name'].apply(fastball_pitch_filter)
    if breaking == True:
        data['pitch_name'] = data['pitch_name'].apply(breaking_pitch_filter)
    if change == True:
        data['pitch_name'] = data['pitch_name'].apply(change_pitch_filter)

    X = data.loc[:, ['release_speed', 'vx0', 'vy0', 'vz0', 'ax', 'ay', 'az', 'release_spin_rate']]
    y = data.loc[:, ['pitch_name']]

    X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

    dt_model = tree.DecisionTreeClassifier(max_depth=10, min_samples_split=50)

    dt_model.fit(X_train, y_train)

    predictions = dt_model.predict(X_train)
    print('Training Score Accuracy {0}'.format(accuracy_score(predictions, y_train)))

    predictions = dt_model.predict(X_test)
    print('Test Score Accuracy {0}'.format(accuracy_score(predictions, y_test)))

    print(classification_report(predictions, y_test))
```

The baseline tree, that is without grouping any pitch, as a training accuracy of 83.6% and test accuracy of 83%. This means that the model correctly labels around 83-84% of all of the 750,000 pitches thrown throughout the 2022 season.

```
In [ ]: baseline_df = decision_tree(pitches, fastball=False, breaking=False, change=False)
```

Training Score Accuracy 0.8357235613776098

Test Score Accuracy 0.8308889362609252

	precision	recall	f1-score	support
4-Seam Fastball	0.94	0.90	0.92	49908
Changeup	0.89	0.85	0.87	17531
Curveball	0.77	0.76	0.76	11627
Cutter	0.54	0.64	0.59	8666
Eephus	0.79	0.87	0.83	99
Knuckle Curve	0.23	0.57	0.33	1261
Knuckleball	0.17	1.00	0.29	1
Sinker	0.82	0.86	0.84	21324
Slider	0.86	0.77	0.81	29394
Split-Finger	0.31	0.74	0.44	919
accuracy			0.83	140730
macro avg	0.63	0.80	0.67	140730
weighted avg	0.85	0.83	0.84	140730

Using just the fastball filter, the model's overall accuracy increases from around 83% to around 89%, with fastball recognition jumping to around 93%. The model does seem to struggle with certain offspeed pitches all around, particularly the knuckle and knuckle curveball, and screwball having an accuracy below 20%, with the knuckleball and the screwball having an accuracy of 0%. This might be due to the variance in how certain pitcher throw these types of pitches, but those will be broken down in the next two scenarios.

```
In [ ]: fastball_df = decision_tree(pitches, fastball=True, breaking=False, change=False)
```

Training Score Accuracy 0.8899642398435293

Test Score Accuracy 0.8858239181411213

	precision	recall	f1-score	support
Changeup	0.89	0.86	0.88	17128
Curveball	0.77	0.76	0.76	11511
Eephus	0.87	0.85	0.86	112
Fastball	0.93	0.95	0.94	81462
Knuckle Curve	0.24	0.60	0.34	1248
Knuckleball	0.17	1.00	0.29	1
Slider	0.86	0.78	0.82	29268
accuracy			0.89	140730
macro avg	0.68	0.83	0.70	140730
weighted avg	0.89	0.89	0.89	140730

Applying the breaking ball filter, the model can surprisingly classify those types of pitches with around 95% accuracy. This could mean that the model might just struggle with individual pitches, or that the broader categorization might mean that the model can label them more accurately.

```
In [ ]: breaking_df = decision_tree(pitches, fastball=False, breaking=True, change=False)
```

Training Score Accuracy 0.8805934089750354

Test Score Accuracy 0.8767711220066795

	precision	recall	f1-score	support
4-Seam Fastball	0.94	0.91	0.92	49474
Breaking	0.95	0.90	0.92	43136
Changeup	0.89	0.85	0.87	17452
Cutter	0.52	0.66	0.58	8119
Eephus	0.70	0.88	0.78	86
Sinker	0.82	0.86	0.84	21517
Split-Finger	0.32	0.73	0.44	946
accuracy			0.88	140730
macro avg	0.73	0.83	0.77	140730
weighted avg	0.89	0.88	0.88	140730

The change grouping has a lower accuracy compared to both the fastball and breaking ball groupings at 90%. It's possible that this is lower due to the fewer number of true "change" pitches that are within the dataset.

```
In [ ]: change_df = decision_tree(pitches, fastball=False, breaking=False, change=True)
```

Training Score Accuracy 0.8362245233311483

Test Score Accuracy 0.831009734953457

	precision	recall	f1-score	support
4-Seam Fastball	0.94	0.91	0.92	49816
Change	0.89	0.85	0.87	17544
Curveball	0.75	0.76	0.76	11269
Cutter	0.52	0.65	0.57	8256
Knuckle Curve	0.23	0.57	0.33	1261
Knuckleball	0.50	1.00	0.67	3
Sinker	0.82	0.86	0.84	21313
Slider	0.87	0.77	0.82	29822
Split-Finger	0.42	0.63	0.51	1446
accuracy			0.83	140730
macro avg	0.66	0.78	0.70	140730
weighted avg	0.85	0.83	0.84	140730

With all three filters active, the model can identify with around 93% accuracy correctly if the next pitch will be either a fastball, breaking ball, or change.

```
In [ ]: combined_df = decision_tree(pitches, fastball=True, breaking=True, change=True)
```

```

Training Score Accuracy 0.9325992997191416
Test Score Accuracy 0.9291835429545939
      precision    recall  f1-score   support

Breaking      0.94      0.91      0.92      42537
Change        0.89      0.86      0.88      17273
Fastball      0.93      0.96      0.94      80920

accuracy              0.93      140730
macro avg      0.92      0.91      0.91      140730
weighted avg   0.93      0.93      0.93      140730

```

Next we will see how well the model does with individual pitchers to see if it can classify their pitches. First up will be Los Angeles Angels superstar Shohei Ohtani, who was picked because is probably the most well-known player in the MLB. For each of these players, we will use pitching data from the last 3 seasons, since single-season data will not be enough to correctly their pitches. As a result, the model can identify Ohtani's pitches with around 97% accuracy, even though he has a fairly large pitching repertoire.

```

In [ ]: ohtani_stats = pyb.statcast_pitcher('2019-03-20', '2022-11-20', 660271)
print("Shohei Ohtani has thrown {} pitches".format(len(ohtani_stats)))
decision_tree(ohtani_stats, fastball=False, breaking=False, change=False)

```

```

Gathering Player Data
Shohei Ohtani has thrown 4938 pitches
Training Score Accuracy 0.9777951933124347
Test Score Accuracy 0.9696969696969697
      precision    recall  f1-score   support

4-Seam Fastball      0.99      0.99      0.99      330
Curveball            0.98      0.95      0.97       66
Cutter                0.91      0.89      0.90      104
Sinker                1.00      0.90      0.95       21
Slider                0.55      0.75      0.63        8
Slow Curve            0.00      0.00      0.00        0
Split-Finger          0.98      0.98      0.98      146
Sweeper               0.97      0.98      0.98      282

accuracy              0.97      957
macro avg      0.80      0.81      0.80      957
weighted avg   0.97      0.97      0.97      957

```

The next player will be New York Yankees ace Gerritt Cole, who is known for a small number of pitches thrown, and his primary pitch is the fastball. Because of the latter fact, we will break down his pitches into grouped fastballs and non-grouped fastballs to compare their accuracies.

```

In [ ]: cole_stats = pyb.statcast_pitcher('2019-03-20', '2022-11-20', 543037)
print("Gerritt Cole has thrown {} pitches".format(len(cole_stats)))

```

```

Gathering Player Data
Gerritt Cole has thrown 12546 pitches

```


It turns out there is not much of a difference between the non-grouped and grouped fastball decision trees, with the non-grouped getting around 98% accurate and the grouped only slightly higher at around 99% accuracy. The model seems to struggle with Cole's sinker, but nearly every other pitch has an accuracy above 80% in the non-grouped model.

```
In [ ]: #Distinct fastballs over 3 season span
decision_tree(cole_stats,fastball=False,breaking=False,change=False)
#Treat 4-seam, 2-seam, and sinker as fastballs
decision_tree(cole_stats,fastball=True,breaking=False,change=False)
```

Training Score Accuracy 0.9859599145386102

Test Score Accuracy 0.985353946297803

	precision	recall	f1-score	support
4-Seam Fastball	0.99	0.99	0.99	1244
Changeup	1.00	1.00	1.00	214
Cutter	0.86	0.82	0.84	39
Knuckle Curve	1.00	1.00	1.00	372
Sinker	0.26	0.36	0.30	14
Slider	0.99	0.99	0.99	575
accuracy			0.99	2458
macro avg	0.85	0.86	0.85	2458
weighted avg	0.99	0.99	0.99	2458

Training Score Accuracy 0.9947095330145488

Test Score Accuracy 0.9914564686737185

	precision	recall	f1-score	support
Changeup	0.99	1.00	0.99	212
Fastball	0.99	0.99	0.99	1293
Knuckle Curve	1.00	1.00	1.00	372
Slider	0.99	0.98	0.99	581
accuracy			0.99	2458
macro avg	0.99	0.99	0.99	2458
weighted avg	0.99	0.99	0.99	2458

The next player analyzed will be San Diego Padres pitcher Yu Darvish, who is known for having a very large pitching repertoire. Without grouping the pitches together, the model can classify Darvish's pitch with around 97% accuracy. It may be how Darvish throw his pitches, but the model does not struggle with the same types of pitches that had issues with both Ohtani's and Cole's pitches.

```
In [ ]: darvish_stats = pyb.statcast_pitcher('2019-03-20', '2022-11-20', 506433)
print("Yu Darvish has thrown {} pitches".format(len(darvish_stats)))
decision_tree(darvish_stats,fastball=False,breaking=False,change=False)
```

Gathering Player Data

Yu Darvish has thrown 10489 pitches

Training Score Accuracy 0.9632417513486388

Test Score Accuracy 0.9473156046161565

	precision	recall	f1-score	support
4-Seam Fastball	0.96	0.95	0.95	497
Changeup	0.00	0.00	0.00	0
Curveball	0.87	0.87	0.87	78
Cutter	0.76	0.82	0.79	78
Knuckle Curve	0.81	0.91	0.85	64
Sinker	0.89	0.88	0.89	144
Slider	0.99	0.97	0.98	614
Slow Curve	0.00	0.00	0.00	0
Split-Finger	0.95	0.96	0.96	151
Sweeper	0.98	0.98	0.98	367
accuracy			0.95	1993
macro avg	0.72	0.73	0.73	1993
weighted avg	0.95	0.95	0.95	1993

The final player being analyzed is Kansas City Royals reliever Aroldis Chapman, who is primarily known for his pitch speed, at one point having the record for the fastest pitch thrown in the majors. Unsurprisingly, the model does have an issue with Chapman's pitches, getting around 99% accurate overall and not less than 88% accurate on any of his 4 pitches.

```
In [ ]: chapman_stats = pyb.statcast_pitcher('2019-03-20', '2022-11-20', 547973)
print("Aroldis Chapman has thrown {0} pitches".format(len(chapman_stats)))
decision_tree(chapman_stats, fastball=False, breaking=False, change=False)
```

Gathering Player Data

Aroldis Chapman has thrown 3204 pitches

Training Score Accuracy 0.9944044764188649

Test Score Accuracy 0.9856230031948882

	precision	recall	f1-score	support
4-Seam Fastball	1.00	0.98	0.99	377
Changeup	0.00	0.00	0.00	1
Sinker	0.88	1.00	0.94	37
Slider	1.00	0.99	0.99	168
Split-Finger	0.91	1.00	0.96	43
accuracy			0.99	626
macro avg	0.76	0.79	0.78	626
weighted avg	0.99	0.99	0.99	626

Part 2: Pitch Prediction

Now that Bob can identify what types of pitches are being thrown, he is ready to step up to the plate and start swinging. Bob has found that he gets better hits on fastballs, so he will train a new model to predict whether or not the next pitch will be a fastball. We will try to find qualitative environmental variables to determine whether or not the next pitch will be a strike or not. These variables should be known before the pitch is thrown, such as contextual information about the game itself. The variables that will be used are the inning, at bat number, pitch number, the number of balls and strikes, the number of outs when the batter is up, and the number of runners on first, second, third, and the total amount of baserunners.

```
In [ ]: features = ['inning', 'at_bat_number', 'pitch_number',  
                  'balls', 'strikes', 'outs_when_up',  
                  'runner_on_first', 'runner_on_second', 'runner_on_third', 'total_baserun  
  
response = ['is_fastball']
```

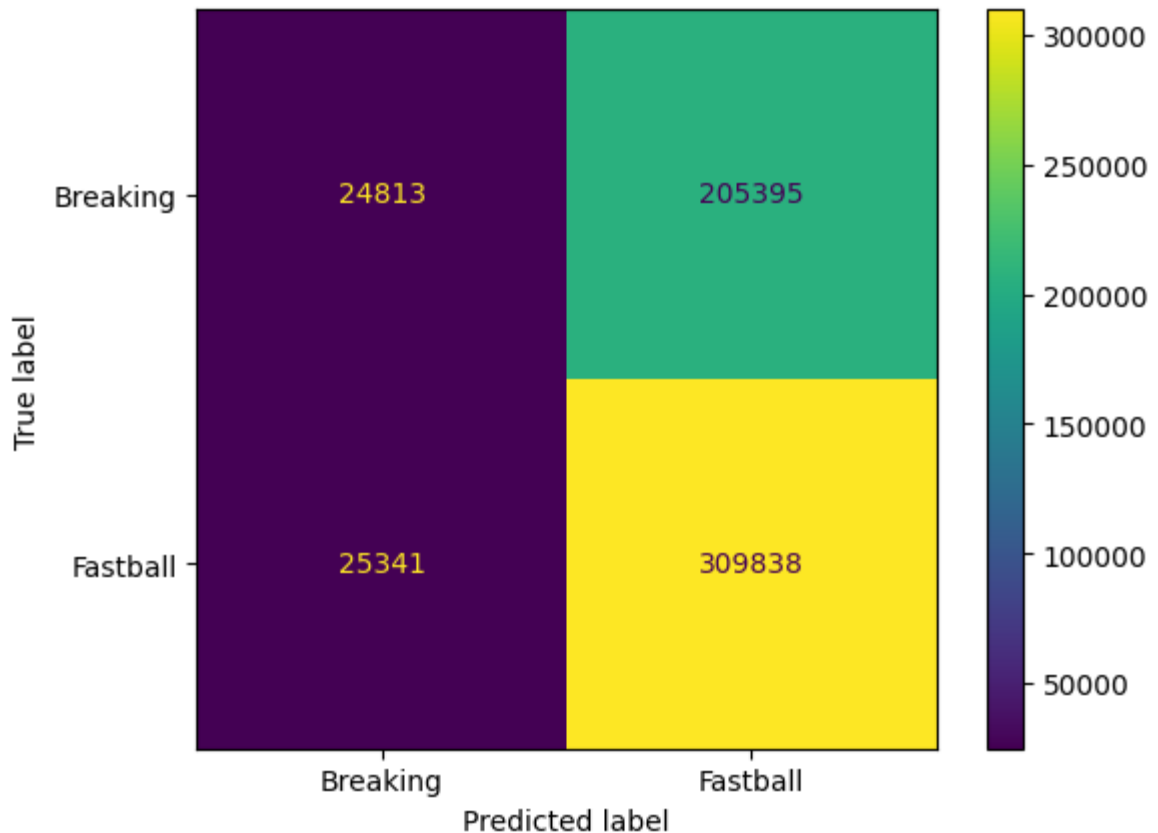
We will use a random forest classification to model our data. Random forests use multiple decision trees and combines their results to reach a single, unified result. First, we will train the data using a random 20% of the pitches dataset that was created at the beginning, then fit that data using the random forest. We will also store the predicted and actual probabilities for the model metrics.

```
In [ ]: forrest = RandomForestClassifier(n_estimators=200, n_jobs=-1, max_depth=10, max_fea  
test, train = train_test_split(pitches, test_size=0.2)  
  
forrest.fit(train[features], train[response])  
  
probabilities = forrest.predict_proba(test[features])  
predictions = forrest.predict(test[features])
```

The AUC of 0.59 means that the model is only slightly better than just blindly picking if the next is a fastball or not. The AUC is also only negligibly better than the null hypothesis, meaning it is doing ever so slightly better than expected. The confusion matrix seems to accurately predict fastballs, but struggles with anything related to offspeed pitches. The false positive rate of around 85% is also mildly concerning, especially with the high amount of offspeed pitches predicted to be fastballs. This could be the result of the ball moving late when it arrives at the plate, which also can confuse the batter.

```
In [ ]: #model performance metrics  
print('AUC: ' + str(metrics.roc_auc_score(y_score=probabilities[:,1], y_true=test[r  
print('Log Loss: ' + str(metrics.log_loss(y_pred=probabilities[:,1], y_true=test[re  
print('Accuracy Score: ' + str(metrics.accuracy_score(test[response], predictions))  
print('Null Score: ' + str(max(test[response].is_fastball.mean(), 1 - test[response  
metrics.ConfusionMatrixDisplay(confusion_matrix=metrics.confusion_matrix(test[respo  
plt.show())
```

```
AUC: 0.5945663453954412  
Log Loss: 0.6570753463959728  
Accuracy Score: 0.5918972314538538  
Null Score: 0.5928311050660875
```



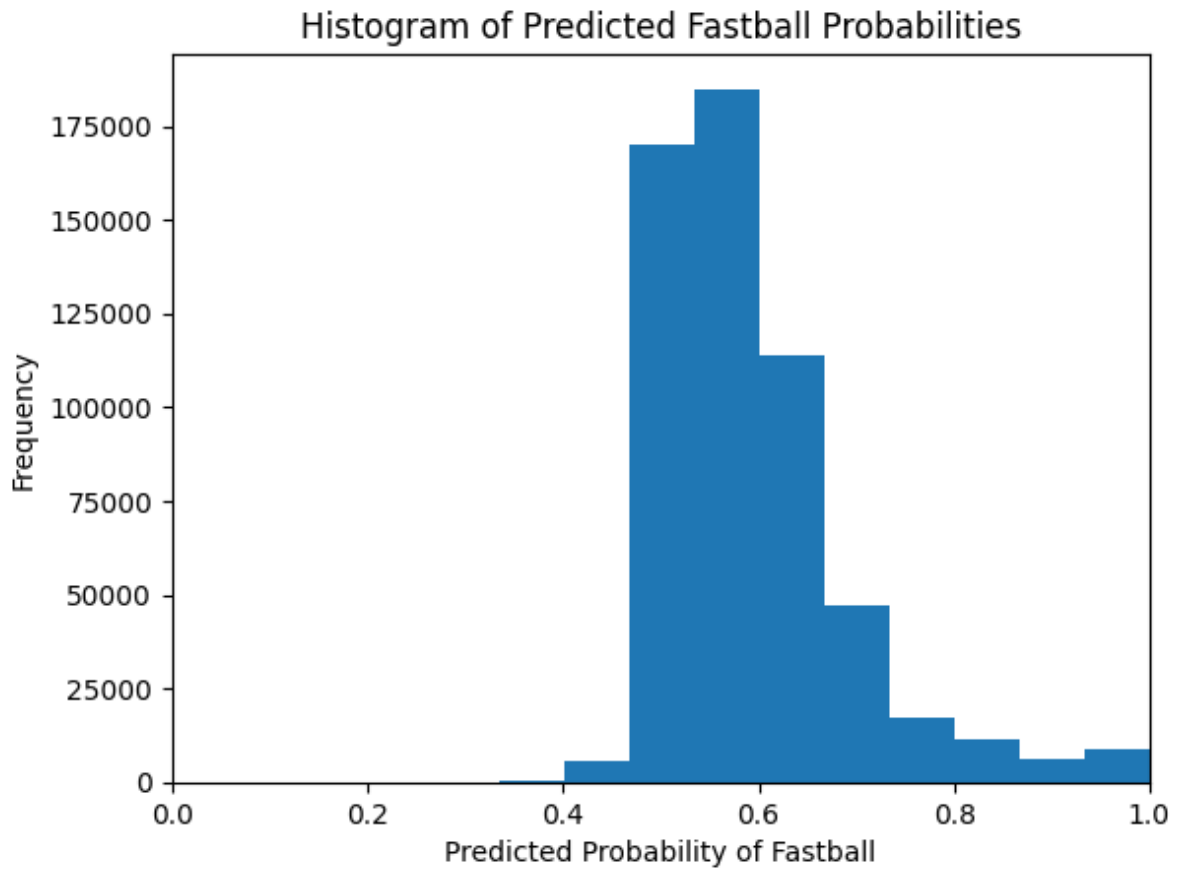
```
In [ ]: confusion = metrics.confusion_matrix(test[response], predictions)
true_pos = confusion[1, 1]
true_neg = confusion[0, 0]
false_pos = confusion[0, 1]
false_neg = confusion[1, 0]

print('Accuracy of CM: ' + str((true_pos+true_neg)/(true_pos+true_neg+false_pos+false_neg)))
print('False Negative Rate of CM: ' + str(false_neg/(false_neg+true_pos)))
print('False Positive Rate of CM: ' + str(false_pos/(false_pos+true_neg)))

Accuracy of CM: 0.5918972314538538
False Negative Rate of CM: 0.07560437855593578
False Positive Rate of CM: 0.8922148665554629
```

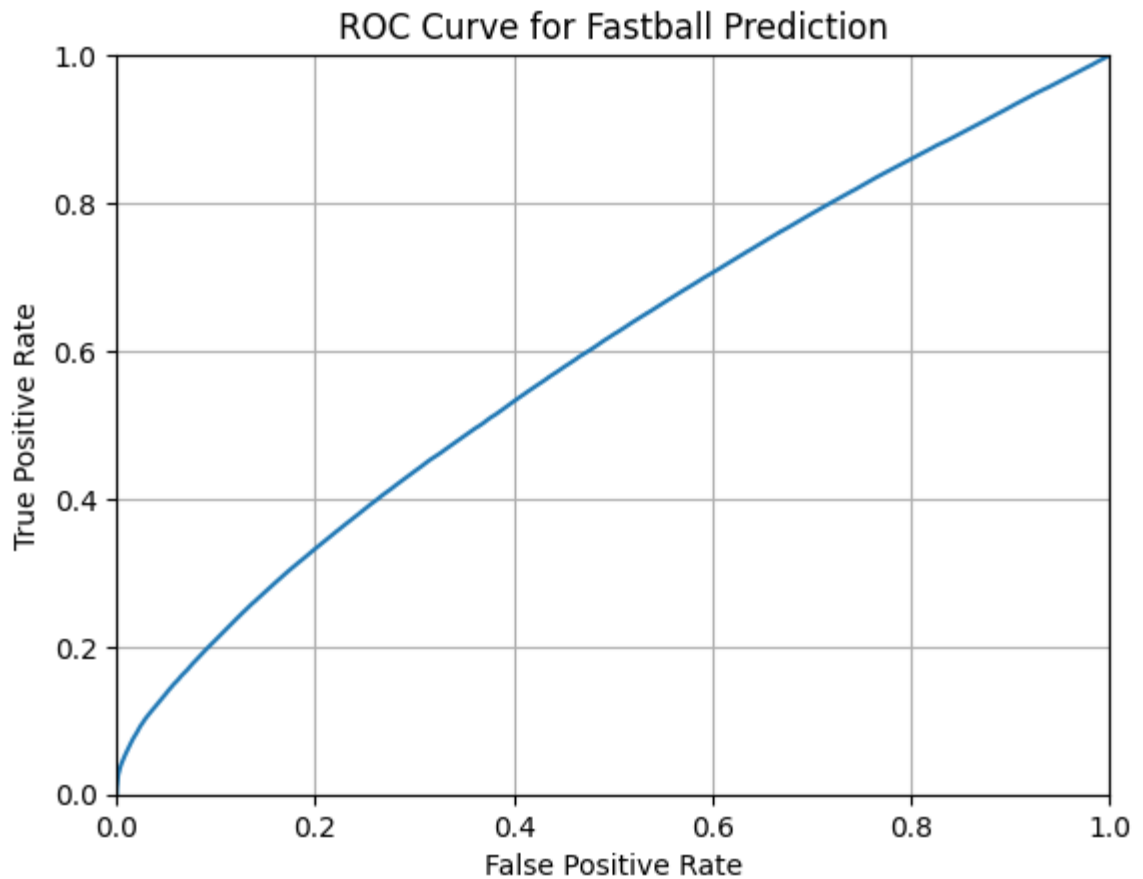
The histogram shows that the model is not very confident in making fastball predictions. There are circumstances where it is very confident, such as a 3-0 count where the pitcher usually throws a fastball to get the batter to swing and keep the at-bat going. Outside of those few circumstances, most of the probabilities are around 45-75% of pitches predicted to be a fastball.

```
In [ ]: plt.hist(probabilities[:,1], bins=10)
plt.xlim(0,1)
plt.title('Histogram of Predicted Fastball Probabilities')
plt.xlabel('Predicted Probability of Fastball')
plt.ylabel('Frequency')
plt.show()
```



The ROC curve does not show a lot of certainty between the 2 types of pitches' certainty, as the curve is trending towards a linear line. As a result, this model needs some work and is most likely underfit.

```
In [ ]: fpr, tpr, thresholds = metrics.roc_curve(test[response], probabilities[:,1])
plt.plot(fpr, tpr)
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.0])
plt.title('ROC Curve for Fastball Prediction')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.grid(True)
plt.show()
```



The best way for Bob and the model to improve is to add more features. This could most likely be achieved by figuring out more environmental variables, such as number of pitches in the at-bat and foul balls hit during the at-bat. Another way to improve the model could be looking at more seasons than just the 2022 season, possibly incorporating data from the 2010-2021 seasons or including them all. However, the hardware might not be able to include all that data, especially if just one season have over 750,000 data points.